



JS Profiling @ TPAC 2021

acomminos@fb.com

cpescheloche@fb.com



Recap



API overview

- A web-exposed sampling profiler for client JS execution
- Provide insight into client performance characteristics on real users' devices
- Shipped in Chrome 94
- github.com/WICG/js-self-profiling

```
const profiler = new Profiler({ sampleInterval: 10, maxBufferSize: 10000 });
```

Starting a profiler

```
const profiler = new Profiler({ sampleInterval: 10, maxBufferSize: 10000 });  
  
window.addEventListener('load', async () => {  
  const trace = await profiler.stop();  
  const traceJson = JSON.stringify(trace);  
  sendTrace(traceJson);  
});
```

Sending trace on load

```
const profiler = new Profiler({ sampleInterval: 10, maxBufferSize: 10000 });

window.addEventListener('load', async () => {
  const trace = await profiler.stop();
  const traceJson = JSON.stringify({
    entries: performance.getEntries(),
    trace,
  });
  sendTrace(traceJson);
});
```

Including data from the performance timeline

```
{  
  "frames": [...],  
  "resources": [...],  
  "samples": [...],  
  "stacks": [...],  
}
```

```
[  
  {  
    "column": 80,  
    "line": 311,  
    "name": "caller",  
    "resourceId": 0  
  },  
  {  
    "column": 368,  
    "line": 311,  
    "name": "callee",  
    "resourceId": 0  
  },  
]
```

Trace format: frames

```
{  
  "frames": [...],  
  "resources": [...],  
  "samples": [...],  
  "stacks": [...],  
}
```



A dotted line originates from the "resources" field in the JSON object and points to a rectangular callout box with a dotted border. The box contains the text `["https://www.fbcdn.net/script.js"]` in green.

```
[ "https://www.fbcdn.net/script.js" ]
```

Trace format: resources

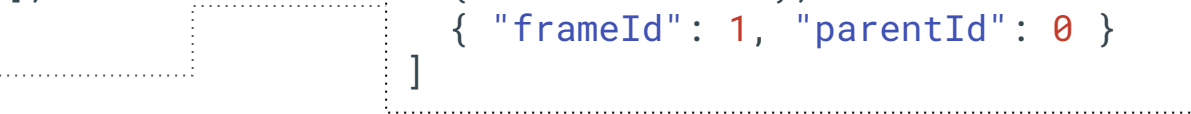

```
{  
  "frames": [...],  
  "resources": [...],  
  "samples": [...],  
  "stacks": [...],  
}
```



```
[  
  {  
    "stackId": 1,  
    "timestamp": 15199  
  },  
  {  
    "timestamp": 15209  
  },  
],
```

Trace format: samples

```
{  
  "frames": [...],  
  "resources": [...],  
  "samples": [...],  
  "stacks": [...],  
}
```



A dotted line connects the "stacks" field in the left JSON object to a callout box on the right. The callout box contains an array of two objects. The first object has "frameId": 0. The second object has "frameId": 1 and "parentId": 0. The numbers 0 and 1 are red, while the strings are blue.

```
[  
  { "frameId": 0 },  
  { "frameId": 1, "parentId": 0 }  
]
```

Trace format: stacks



What's working well?

- Initial data suggests **enabling profiling slows load time by <1% ($p=0.05$)** at FB
 - Strong evidence that sampling profiling can be implemented with minimal overhead
- API has provided a drop-in solution for FB app client perf analysis
- Strong adoption from other industry partners



What could be better?

- Non-JS execution is hard to identify in traces
 - Currently, top-level UA work is indistinguishable from idle execution
 - GC activity adds to the noise of long traces
 - Client code that causes asynchronous rendering work isn't measurable
- Interactions with Long Tasks API can be cumbersome

Representing non-JS execution



Introducing state markers

- Tags a sample with a string representing top level UA work category
- Similar representation to traces visualized through devtool profilers



Marker candidates

- Need to be generic and interoperable
- Script related:
 - script: js execution, optional ?
 - parse: HTML? JS?
 - gc
- Rendering related:
 - paint: update the rendering part of the event loop or limited to actual paint
 - style
 - layout



API Modification

```
enum ProfilerMarker { "script", "gc", "parse", "paint", "other" };  
  
dictionary ProfilerSample {  
    required DOMHighResTimeStamp timestamp;  
    unsigned long stackId;  
    ProfilerMarker? marker;  
};
```



```
"samples" : [  
  {  
    "timestamp" : 100,  
    "stackId": 3  
  },  
  {  
    "timestamp" : 110,  
    "stackId": 2  
  },  
  {  
    "timestamp" : 120,  
    "stackId": 2  
  },  
  {  
    "timestamp" : 130,  
    "stackId": 2  
  },  
  {  
    "timestamp" : 140,  
    "stackId": 1  
  },  
  {  
    "timestamp" : 150  
  }  
]
```

```
"samples" : [  
  {  
    "timestamp" : 100,  
    "stackId": 3,  
    "marker": "script"  
  },  
  {  
    "timestamp" : 110,  
    "stackId": 2,  
    "marker": "script"  
  },  
  {  
    "timestamp" : 120,  
    "stackId": 2,  
    "marker": "gc"  
  },  
  {  
    "timestamp" : 130,  
    "stackId": 2,  
    "marker": "gc"  
  },  
  {  
    "timestamp" : 140,  
    "stackId": 1,  
    "marker": "script"  
  },  
  {  
    "timestamp" : 150  
  }  
]
```

Example trace GC



Security and privacy concerns

- Profiles **must not** expose work done on a cross-origin document
 - Top level UA work may only appear in a trace if the **responsible document** for the work is same-origin with the associated Profiler
- New information exposed, need to limit granularity of marker types
 - Need to avoid introducing new side channels
 - Require cross-origin isolation for markers?



Prototype

- TAG review: <https://github.com/w3ctag/design-reviews/issues/682>
- Intent to prototype:
 - Implementation for review in V8/Blink
 - Target Origin Trial in Chrome 98



Open questions

- Interest in breaking down paint marker into:
 - Style
 - Layout
 - Paint
- Events like GC may be hard to isolate timing by origin
- Is JS self profiling the best place for this information?
 - Performance-timeline could be a candidate

Profiling and long tasks



Issues with the Long Tasks API

- Currently hard to debug root cause for long tasks
 - Partially solved by [Long Tasks V2 API sketch](#), though not much movement
- Can profiling help here?



Supplementing with JS Self-Profiling

- If running a profiler, you can **cross-correlate with recorded samples to root cause**
 - Find expensive sampled functions
 - Identify UA-level work (e.g. GC/paint/layout) with the marker extension



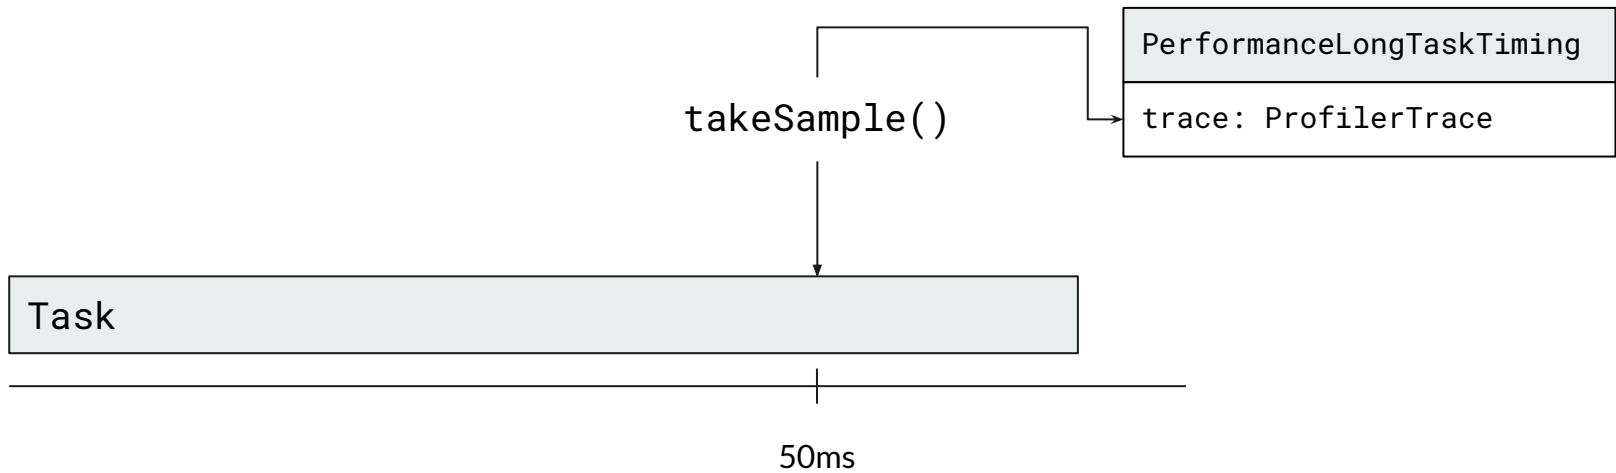
Drawbacks of JSSP correlation

- Requires recording **all samples**, when we may only care about samples responsible for a long task
 - Increased memory and CPU pressure from sampling and trace processing



Can we do better?

- What if we added an option to **sample only during long tasks**?
 - Existing Document-Policy header `js-profiling` can hint to starting maintaining code map metadata
 - Active profiler not necessary
- Requires proactively detecting long tasks
 - Schedule 50ms background task to capture a sample (RAIL long task definition)



Sampling long tasks



API modification

```
[Exposed=Window]
interface PerformanceLongTaskTiming : PerformanceEntry {
    readonly attribute FrozenArray<TaskAttributionTiming> attribution;
    readonly attribute ProfilerTrace? trace;
    [Default] object toJSON();
};
```



Security and privacy

- Need to ensure we do not expose attributions across origins
- Actual sampling is a subset of existing profiling functionality
 - Existing cross-origin checks will continue to make this safe



Open questions

- How should we activate this?
 - Add samples when js-profiler document policy is present?
- Is a single sample enough to get signal?
- Will proactive long task sampling be *slower* than a long-running profiler?
 - Additional main-thread work is likely required before each task runs

Discussion

Appendix



Links

- Explainer: <https://github.com/WICG/js-self-profiling/pull/55>
- Markers TAG review: <https://github.com/w3ctag/design-reviews/issues/682>


```
[Exposed=Window]
interface Profiler : EventTarget {
    readonly attribute DOMHighResTimeStamp sampleInterval;
    readonly attribute boolean stopped;

    constructor(ProfilerInitOptions options);
    Promise<ProfilerTrace> stop();
};
```

IDL: Profiler